




# Chapter 2

Regular Expressions and Word Tokenization

# Outline

- Introduction to regular expressions (**regex**)
- Which pattern?
- Practicing regular expressions: **re.split()** and **re.findall()**
- Introduction to tokenization
- Word tokenization with **NLTK**
- More regex with **re.search()**
- Advanced tokenization with NLTK and regex
- Choosing a tokenizer
- Regex with NLTK tokenization
- Non-ascii tokenization
- Charting word length with NLTK
- Charting practice



# Introduction to regular expressions

# What are regular expressions?

- Regular expressions (**Regex**) are strings with a **special syntax**.
- Regex allow us to **match** patterns in other strings.
- **Applications** of Regex
  - Find all web links in a document
  - Parse email addresses
  - Remove/Replace unwanted characters

```
import re  
re.match('abc', 'abcdef')
```

```
<re.Match object; span=(0, 3), match='abc'>
```

```
word_regex = '\w+'  
re.match(word_regex, 'hi there!')
```

```
<re.Match object; span=(0, 2), match='hi'>
```

# Common regex patterns

- A **pattern** is a series of letters/symbols which can **map** to an actual text/words/punctuation.
- Examples

Pattern	Matches	Example
\w+	Word	'Magic'
\d	Digit	9
\s	Space	' '
\S	Not space	'no_spaces'
[a-z]	Lowercase group	'abc'

- The wildcard will match ANY letter or symbol
- The + and \* characters allow things to become greedy, grabbing repeats of single letters or whole patterns.

More information: <https://docs.python.org/3/library/re.html>

# re module


- re module
  - split: split a string on regex
  - findall: find all patterns in a string
  - search: search for a pattern
  - match: match an entire string or substring based on a pattern

```
re.split('\s+', 'Split on spaces.')
```

```
['Split', 'on', 'spaces.']
```

```
word_regex2 = 'spaces!'
print(re.findall(word_regex2, 'Split on spaces!'))
print(re.search(word_regex2, 'Split on spaces!'))
print(re.match(word_regex2, 'Split on spaces!'))
```

```
['spaces']
<re.Match object; span=(9, 15), match='spaces'>
None
```



Let's practice!

# Which pattern?

- Which of the following Regex patterns results in the following text?

```
my_string = "Let's write RegEx!"  
re.findall(PATTERN, my_string)
```

```
['Let', 's', 'write', 'RegEx']
```

- Possible Answers

- a) PATTERN = r"\s+"
- b) PATTERN = r"\w+"
- c) PATTERN = r"[a-z]"
- d) PATTERN = r"\w"



# re.split() and re.findall()

- The regular expression module `re` which is imported from you
- Set `my_string = "Let's write RegEx! Won't that be fun? I sure think so. Can you find 4 sentences? Or perhaps, all 19 words?"`
- Note: It's important to prefix your regex patterns with `r` to ensure that your patterns are interpreted in the way you want them to. Else, you may encounter problems to do with **escape sequences** in strings.
  - For example, `"\n"` in Python is used to indicate a new line, but if you use the `r` prefix, it will be interpreted as the raw string `"\n"` - that is, the character `"\"` followed by the character `"n"` - and not as a new line.

# re.split() and re.findall()

- Split **my\_string** on each sentence ending. To do this:
  - Write a pattern called **sentence\_endings** to match sentence endings ( **.?!**  ).
  - Use **re.split()** to split **my\_string** on the pattern and print the result.

```
# Write a pattern to match sentence endings: sentence_endings
sentence_endings = r"["
```

```
# Split my_string on sentence endings and print the result
print(re.__(__, __))
```

```
["Let's write RegEx", " Won't that be fun", ' I sure think so', ' Can you find 4 sentences',  
' Or perhaps, all 19 words', ""]
```

# re.split() and re.findall()

- Find and print all capitalized words in `my_string` by writing a pattern called `capitalized_words` and using `re.findall()`.
  - Remember the `[a-z]` pattern to match lowercase groups? Modify that pattern appropriately in order to match uppercase groups.

```
# Find all capitalized words in my_string and print the result
capitalized_words = r"__\w+"
print(re.__(__, __))
```

```
['Let', 'RegEx', 'Won', 'Can', 'Or']
```

# re.split() and re.findall()

- Write a pattern called **spaces** to match one or more spaces ("**\s+**") and then use **re.split()** to split **my\_string** on this pattern, keeping all punctuation intact. Print the result.

```
# Split my_string on spaces and print the result
spaces = r"_____"
print(re. ____ (____, ____))
```

```
["Let's", 'write', 'RegEx!', "Won't", 'that', 'be', 'fun?', 'I', 'sure', 'think', 'so.', 'Can',  
'you', 'find', '4', 'sentences?', 'Or', 'perhaps,', 'all', '19', 'words?']
```

# re.split() and re.findall()


- Find all digits in `my_string` by writing a pattern called digits (`"\d+"`) and using `re.findall()`. Print the result.

```
# Find all digits in my_string and print the result
digits = r"_____"
print(re.findall(____, ____))
```

```
['4', '19']
```

## Tokenization

Natural Language Processing



[ 'Natural', 'Language', 'Processing' ]

# Introduction to tokenization

# What is tokenization?

- Turning a string or document into **tokens** (smaller chunks).
- One step for preprocessing a text in NLP
- Many different theories and rules of tokenization, so, everyone can create their rules using Regex.
- Examples:
  - break out words or sentences
  - separate punctuation
  - separating all hashtags in a Tweet

# nltk library

- nltk: natural language toolkit
  - `py -m pip install nltk`
- example of using the **word\_tokenize** method to break down a string into tokens.

```
from nltk.tokenize import word_tokenize  
word_tokenize("Hi there!")
```

```
['Hi', 'there', '!']
```



# Why tokenize?

- Easier to map part of speech
- Matching common words
- Removing unwanted tokens
- Example:
  - The sentence is: "I don't like Sam's shoes."
  - Result after tokenization is:


```
['I', 'do', "n't", 'like', 'Sam', "'s", 'shoes', '.']
```

# Other nltk tokenizers

- **sent\_tokenize**: tokenize a document into sentences
- **regexp\_tokenize**: tokenize a string or document based on a regular expression pattern
- **TweetTokenizer**: special class just for tweet tokenization, allowing you to separate hashtags, mentions and lots of exclamation points (!!!).

## Tokenization

Natural Language Processing



[ 'Natural', 'Language', 'Processing' ]

Let's practice!

# Word tokenization with NLTK

- Utilize `word_tokenize` and `sent_tokenize` from `nltk.tokenize` to tokenize both words and sentences from Python strings.
- Import the `sent_tokenize` and `word_tokenize` functions from `nltk.tokenize`  
`# Import necessary modules`  
`from nltk.tokenize import sent_tokenize`  
`from nltk.tokenize import word_tokenize`
- Read txt file and keep in `scene_one`  
`#Read TXT file`  
`f = open("scene_one.txt", "r")`  
`scene_one = f.read()`

# Word tokenization with NLTK

- Tokenize all sentences in `scene_one` using the `sent_tokenize()` function.

```
# Split scene_one into sentences: sentences
sentences = _____(_____)
```

- Tokenize the fourth sentence in `sentences`, which you can access as `sentences[3]`, using the `word_tokenize()` function.

```
# Use word_tokenize to tokenize the fourth sentence:
tokenized_sent
tokenized_sent = _____(_____)
```

# Word tokenization with NLTK

- Find the unique tokens in the entire scene by using `word_tokenize()` on `scene_one` and then converting it into a set using `set()`.

```
# Make a set of unique tokens in the entire scene: unique_tokens  
unique_tokens = _____(_____())
```

- Print the unique tokens found.

```
# Print the unique tokens result  
print(unique_tokens)
```

# More regex with re.search()

- Utilize `re.search()` and `re.match()` to find specific tokens.
- Use `re.search()` to search for the first occurrence of the word "coconuts" in `scene_one`. Store the result in `match`.

```
# Search for the first occurrence of "coconuts" in  
scene_one: match
```

```
match = re._____(_____, scene_one)
```

- Print the start and end indexes of `match` using its `.start()` and `.end()` methods, respectively.

```
# Print the start and end indexes of match  
print(_____, _____)
```

# More regex with re.search()

- Write a regular expression called **pattern1** to find anything in square brackets.

```
# Write a regular expression to search for anything in  
square brackets: pattern1  
pattern1 = r"\[____\]"
```

- Use **re.search()** with the pattern to find the first text in **scene\_one** in square brackets in the scene. Print the result.

```
# Use re.search to find the first text in square brackets  
print(re._____(_____, _____))
```




# More regex with re.search()

- Create a pattern to match the script notation (e.g. **Character:**), assigning the result to **pattern2**. Remember that you will want to match any words or spaces that precede the : (such as the space within **SOLDIER #1:**).
- Use **re.match()** with your new pattern to find and print the script notation in the fourth line. The tokenized sentences are available in your namespace as **sentences**.

```
# Find the script notation at the beginning of the fourth
sentence and print it
pattern2 = r"[_]+:"
print(re.__(__(, __[3])))
```

## Tokenization

Natural Language Processing



[ 'Natural', 'Language', 'Processing' ]

Advanced tokenization with NLTK and regex

# Regex groups using or "|"

- OR method is represented using |
- Define a group using ()
- Define explicit character ranges using []

```
import re
match_digits_and_words = ('(\d+|\w+)')
re.findall(match_digits_and_words, 'He has 11 cats')
```

```
['He', 'has', '11', 'cats']
```

# Regex ranges and groups

- Examples

Pattern	Matches	Example
[A-Za-z]+	Upper and lowercase English alphabet	'ABCDEFGHghijk'
[0-9]	Number from 0-9	9
[A-Za-z\-\.\.]+	Upper and lowercase English alphabet, - and .	'My-Website.com'
(a-z)	a, - and z	'a-z'
(\s+ ,)	Spaces or a comma	','


# Character range with `re.match()`

```
import re
my_str = 'match lowercase spaces nums like 12, but no commas'
match_str = ('[a-z0-9 ]+')
print(re.match(match_str, my_str))
```

```
<re.Match object; span=(0, 35), match='match lowercase spaces nums like 12'>
```

## Tokenization

Natural Language Processing



[ 'Natural', 'Language', 'Processing' ]

Let's practice!

# Choosing a tokenizer

```
['SOLDIER', '#1', 'Found', 'them', '?', 'In', 'Mercea', '?', 'The',  
'coconut', 's', 'tropical', '!']
```

```
my_string = "SOLDIER #1: Found them? In Mercea? The coconut's tropical!"
```

- Given the following string, which of the below patterns is the best tokenizer? If possible, you want to retain sentence punctuation as separate tokens, but have '#1' remain a single token.
- The string is available in your workspace as **my\_string**, and the patterns have been pre-loaded as **pattern1**, **pattern2**, **pattern3**, and **pattern4**, respectively.
- Additionally, **regex\_tokenize** has been imported from **nlTK.tokenize**. You can use **regex\_tokenize(string, pattern)** with **my\_string** and one of the patterns as arguments to experiment for yourself and see which is the best tokenizer.

## Possible Answers

- a) `r"(\w+|\?|!)"`
- b) `r"(\w+|#\d|\?|!)"`
- c) `r"(\#\d\w+\?|!)"`
- d) `r"\s+"`

# Regex with NLTK tokenization

- Twitter is a frequently used source for NLP text and tasks. In this exercise, you'll build a more complex tokenizer for tweets with hashtags and mentions using `nltk` and `regex`. The `nltk.tokenize.TweetTokenizer` class gives you some extra methods and attributes for parsing tweets.
- From `nltk.tokenize`, import `regex_tokenize` and `TweetTokenizer` .

```
# Import the necessary modules
```

```
from nltk.tokenize import _____
```

```
from nltk.tokenize import _____
```



# Regex with NLTK tokenization

- Set `tweets = ['This is the #nlp exercise! #python', '#NLP is super fun! <3 #learning', 'Thanks @fitmkmtnb :) #nlp #python']`
- A regex pattern to define hashtags called **pattern1** has been defined for you. Call **regex\_tokenize()** with this hashtag pattern on the first tweet in **tweets** and assign the result to **hashtags**.

- Print **hashtags**  
# Define a regex pattern to find hashtags: pattern1  
`pattern1 = r"#\w+"`  
# Use the pattern on the first tweet in the tweets list  
`hashtags = _____(_____[_], _____)`  
`print(hashtags)`

```
['#nlp', '#python']
```

# Regex with NLTK tokenization

- Write a new pattern called **pattern2** to match mentions and hashtags. A mention is something like **@fitmkmtnb**.
- Then, call **regex\_tokenize()** with your new hashtag pattern on the **last** tweet in **tweets** and assign the result to **mentions\_hashtags**.
  - You can access the last element of a list using **-1** as the index, for example, **tweets[-1]**.
- Print **mentions\_hashtags**

```
# Write a pattern that matches both mentions (@) and hashtags
pattern2 = r"([_]\w+)"
# Use the pattern on the last tweet in the tweets list
mentions_hashtags = _____(_____[_], _____)
print(mentions_hashtags)
```

```
['@fitmkmtnb', '#nlp', '#python']
```

# Regex with NLTK tokenization

- Create an instance of **TweetTokenizer** called **tknzs** and use it inside a list comprehension to tokenize each tweet into a new list called **all\_tokens**.
- To do this, use the **.tokenize()** method of **tknzs**, with **t** as your iterator variable.
- Print **all\_tokens** .

```
# Use the TweetTokenizer to tokenize all tweets into one list
tknzs = _____()
all_tokens = [_____._____(_) for t in tweets]
print(all_tokens)
```

```
[['This', 'is', 'the', '#nlp', 'exercise', '!', '#python'], ['#NLP', 'is', 'super', 'fun', '!', '<3', '#learning'], ['Thanks', '@fitmkmtnb', ':)', '#nlp', '#python']]
```

# Non-ascii tokenization

- Practice advanced tokenization by tokenizing some non-ascii based text.
- The following modules have been pre-imported from **nltk.tokenize**: **regex\_tokenize** and **word\_tokenize**.

```
from nltk.tokenize import regex_tokenize
from nltk.tokenize import word_tokenize
```

- Set `german_text = "Wann gehen wir Pizza essen? 🍕 Und fährst du mit Über? 🚗"`
- Unicode ranges for emoji are:
  - ('\\U0001F300'-'\\U0001F5FF'), ('\\U0001F600'-'\\U0001F64F'), ('\\U0001F680'-'\\U0001F6FF'), and ('\\u2600'-'\\u26FF'-'\\u2700'-'\\u27BF').

# Non-ascii tokenization

- Tokenize all the words in `german_text` using `word_tokenize()`, and print the result.

```
# Tokenize and print all words in german_text  
all_words = _____  
print(all_words)
```

```
['Wann', 'gehen', 'wir', 'Pizza', 'essen', '?', '🍕', 'Und', 'fährst', 'du',  
'mit', 'Über', '?', '🚗']
```

# Non-ascii tokenization

- Tokenize only the capital words in `german_text` .
  - First, write a pattern called `capital_words` to match only capital words. Make sure to check for the German `Ü` ! To use this character in the exercise, copy and paste it from these instructions.
  - Then, tokenize it using `regex_tokenize()` .

```
# Tokenize and print only capital words
capital_words = r"[_]\w+"
print(_____(_____, _____))
```

```
['Wann', 'Pizza', 'Und', 'Über']
```

# Non-ascii tokenization


- Tokenize only the emoji in **german\_text**. The pattern using the unicode ranges for emoji given in the assignment text has been written for you. Your job is to use **regex\_tokenize()** to tokenize the emoji.

```
# Tokenize and print only emoji
```

```
emoji = "[ '\U0001F300-\U0001F5FF' | '\U0001F600-  
\U0001F64F' | '\U0001F680-\U0001F6FF' | '\u2600-  
\u26FF\u2700-\u27BF' ]"
```

```
print(_____(_____, _____))
```





Charting word length with nltk



# Getting started with **matplotlib**

- Matplotlib is a charting library used by many open-source Python projects
- Straightforward functionality with lots of options
  - histograms,
  - bar charts,
  - line charts and
  - scatter plots.
- And also advanced functionality like generating 3D graphs and animations.


# Plotting a histogram with matplotlib

Install **matplotlib**:

```
pip install matplotlib
```

```
from matplotlib import pyplot as plt  
plt.hist([1, 5, 5, 7, 7, 7, 9])
```

```
plt.show()
```




Generated histogram

# Combining NLP data extraction with plotting


```
from matplotlib import pyplot as plt
from nltk.tokenize import word_tokenize

words = word_tokenize("This is a pretty
cool tool!")
word_lengths = [len(w) for w in words]
plt.hist(word_lengths)

plt.show()
```



Word length histogram



Let's practice!

# Charting practice

- Find and chart the number of words per line in the script using **matplotlib**.
- Read txt file (holy\_grail.txt) and keep in **holy\_grail**

```
#Read TXT file  
f = open("holy_grail.txt", "r")  
holy_grail = f.read()
```

- Split the script **holy\_grail** into lines using the newline ('\n') character.

```
# Split the script into lines: lines  
lines = _____._____('\n')
```

# Charting practice

- Use `re.sub()` inside a list comprehension to replace the prompts such as **ARTHUR:** and **SOLDIER #1**. The pattern has been written for you.
- Use a list comprehension to tokenize `lines` with `regex_tokenize()`, keeping only words. Recall that the pattern for words is `"\w+"`.

```
# Replace all script lines for speaker
```

```
pattern = "[A-Z]{2,}(\s)?(#\d)?([A-Z]{2,})?:"
```

```
lines = [re.____(____, ' ', 1) for l in lines]
```

```
# Tokenize each line: tokenized_lines
```

```
tokenized_lines = [____(____,____) for s in lines]
```

# Charting practice

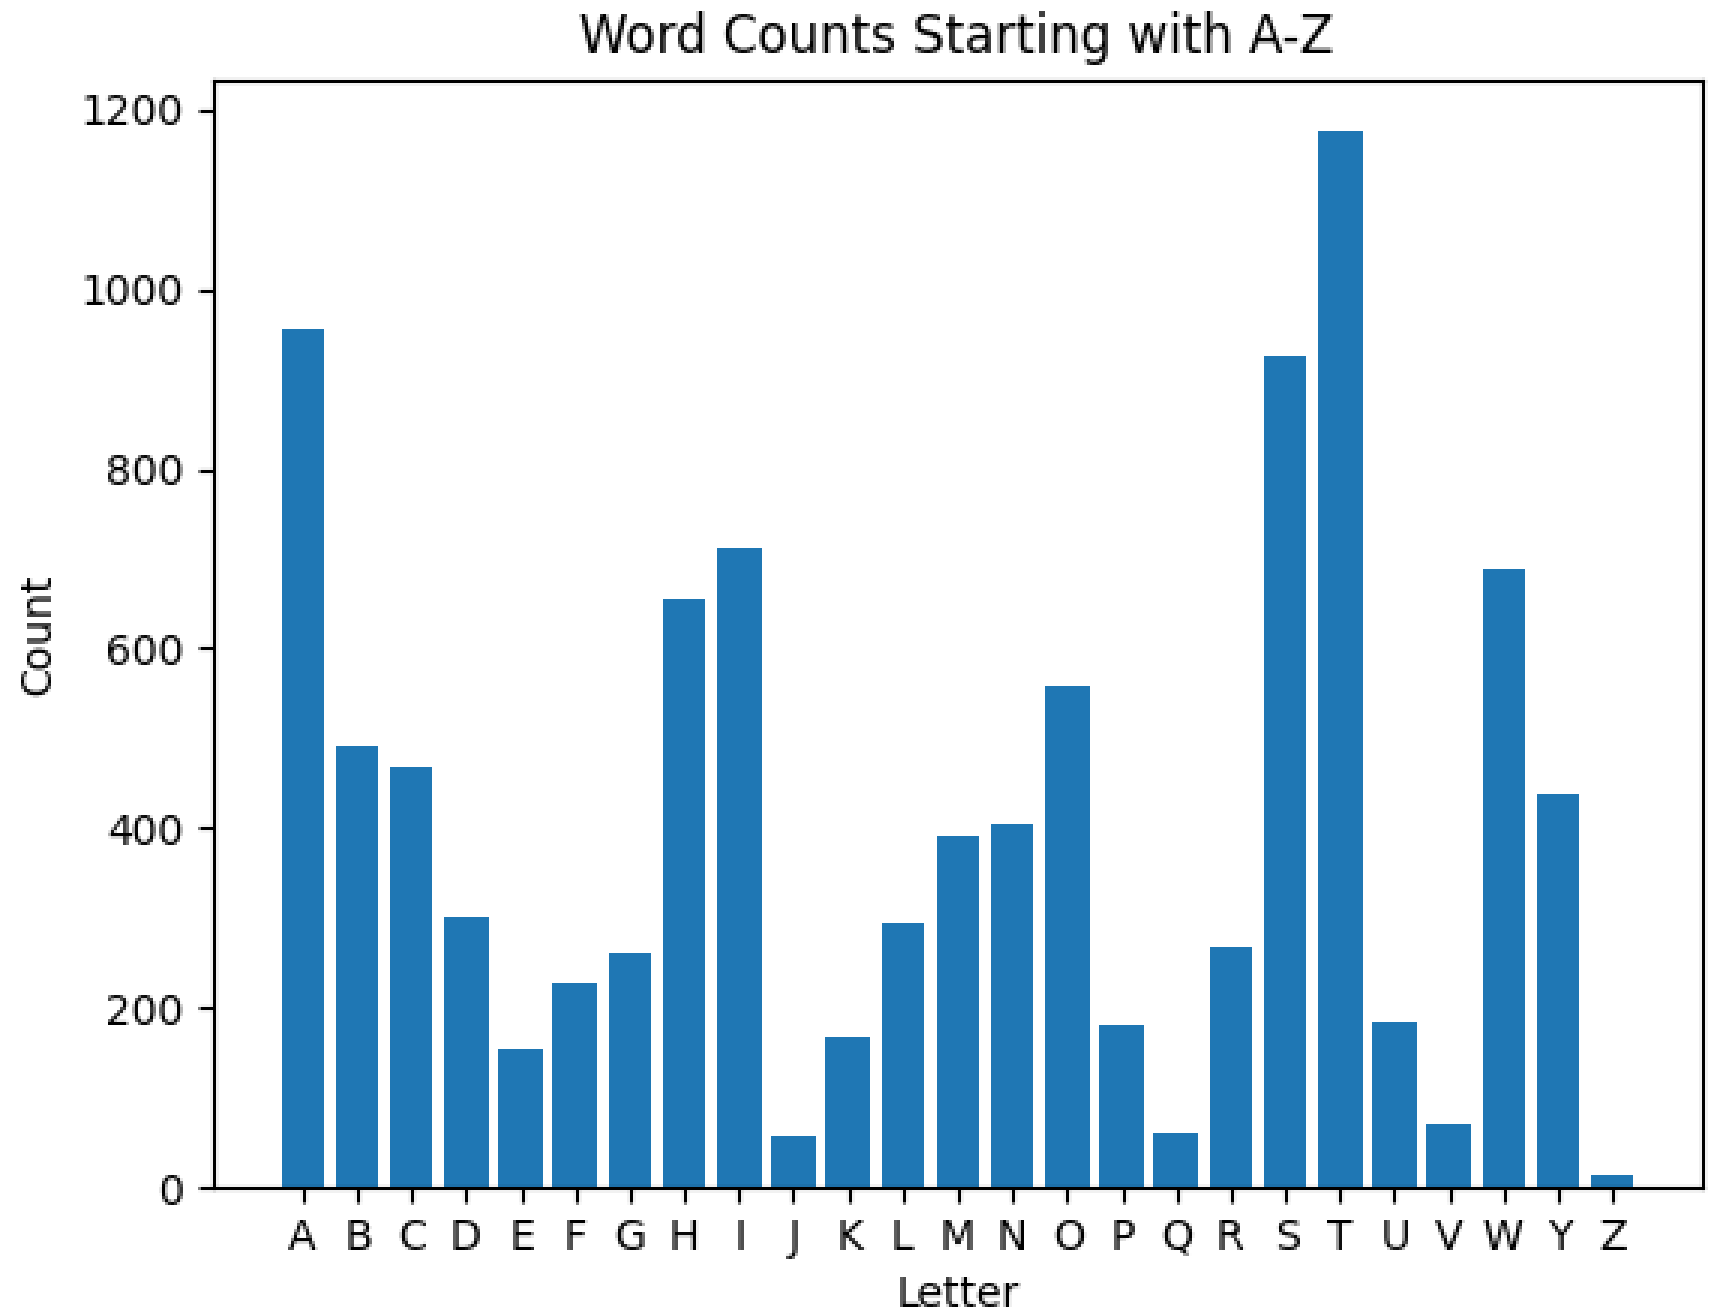
- Use a list comprehension to create a list of line lengths called `line_num_words`.
  - Use `t_line` as your iterator variable to iterate over `tokenized_lines`, and then `len()` function to compute line lengths.

```
# Make a frequency list of lengths: line_num_words
line_num_words = [_____ (_____) for t_line in tokenized_lines]
```

- Plot a histogram of `line_num_words` using `plt.hist()`. Don't forget to use `plt.show()` as well to display the plot.

```
# Plot a histogram of the line lengths
_____._____ (_____)
# Show the plot
_____._____ ()
```

# Work







# Questions

Reference: <https://app.datacamp.com/learn>